



**NOVA**

**IMS**

Information  
Management  
School

# Reinforcement Learning

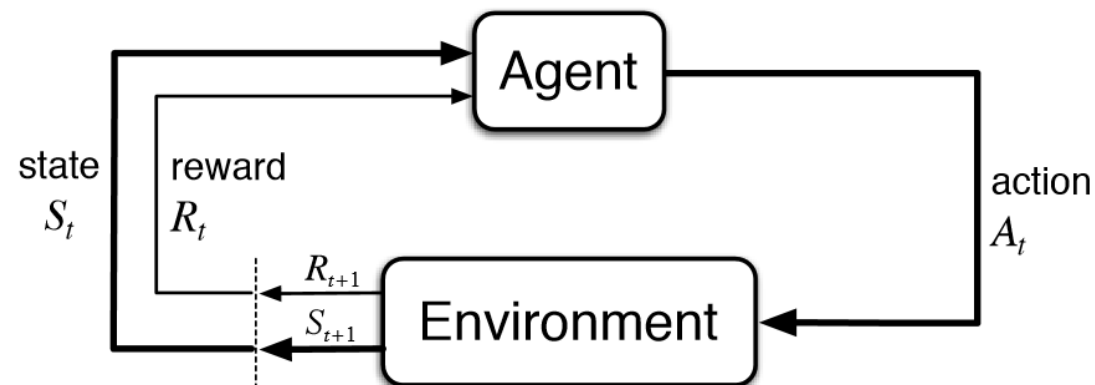
Lecture 4

Model Free Methods

May 2025

Nuno Alpalhão  
nalpalhao@novaims.unl.pt

- All components are functions
  - Policies  $\pi : S \rightarrow A$  (or to probabilities over  $A$ )
  - Value functions  $v : S \rightarrow R$ 
    - Q functions  $q : S \times A \rightarrow R$
  - Models  $m : S \rightarrow S$  and/or  $r : S \rightarrow R$
  - State update  $u : S \times O \rightarrow S$



**Monte Carlo** (MC) is one of the most popular and most commonly used algorithms in various fields ranging from physics and mechanics to computer science. The Monte Carlo algorithm is used in reinforcement learning (RL) when the model of the environment is not known.

Using **Dynamic Programming** (DP) to **find an optimal policy** where **we know the model dynamics**, which is **transition and reward probabilities**. But how can we determine the optimal policy when we don't know the model dynamics? In that case, we use the Monte Carlo algorithm; it is extremely powerful for finding optimal policies when we don't have knowledge of the environment.

In the **Monte Carlo Prediction** setup (in the prediction vs control realm) we are solely trying to achieve a **value function**.

We are **assuming an a priori policy** (which we normally do not have).

## Disadvantages of Monte Carlo Learning

Monte Carlo algorithms can be used to learn value predictions:

- When episodes are long, **learning can be slow.**
- **Wait until an episode ends** before we can learn.
- Return can have **high variance.**

Estimating  $v_\pi$  or  $Q^\pi$  is called policy evaluation or, simply, **prediction**.

- Given a policy, what is my expected return under that behaviour?
- Given this treatment protocol/trading strategy, what is my expected return?

Estimating  $v_*$  or  $Q^*$  is sometimes called **control**, because these can be used for policy optimisation.

- What is the optimal way of behaving? What is the optimal value function?
- What is the optimal treatment? What is the optimal control policy to minimise time, fuel consumption, etc?

Update  $v(s)$  **incrementally** after each episode:

For each state  $S_t$  with return  $G_t$ :

- $N'(S_t) \leftarrow N(S_t) + 1$
- $v'(S_t) \leftarrow v(S_t) + \frac{1}{N(S_t)} (G_t - v(S_t))$

In non-stationary problems, it can be useful to track a running mean, i.e. forget old episodes:

$$v'(S_t) \leftarrow v(S_t) + \alpha (G_t - v(S_t))$$

Learn  $v_\pi$  online from experience under policy  $\pi$ .

Incremental every-visit Monte Carlo:

- Update value  $v(S_t)$  toward actual return  $G_t$ :

$$v'(S_t) \leftarrow v(S_t) + \alpha (G_t - v(S_t))$$

$\alpha$  can take many values:

$$\alpha = \frac{1}{N(S_t)}$$

or even  $\alpha = 1$



Generally high variance estimator:

- Reducing variance can require a lot of data.
- In cases where data is very hard or expensive to acquire, or the stakes are high, Monte Carlo may be impractical.

Requires episodic settings:

- Episode must end before data can be used to update the value function.

“If one had to identify one idea as central and novel to reinforcement learning, it would undoubtedly be temporal-difference (TD) learning.” —  
Sutton and Barto 2017

Temporal Difference methods **learn immediately** from each step in an **episode**.

Temporal Difference **is model-free**: no knowledge of MDP transitions / rewards.

Temporal Difference learns from incomplete episodes, by **bootstrapping**.

Temporal Difference can be used in episodic or **infinite-horizon** non-episodic settings.

Bellman's equations:

$$v_{\pi}(s) = E [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t \sim \pi(S_t)]$$

Approximate by iterating:

$$v'(s) = E [R_{t+1} + \gamma v(S_{t+1}) | S_t = s, A_t \sim \pi(S_t)]$$

We can sample this:

$$v'(S_t) = R_{t+1} + \gamma v(S_{t+1})$$

We can sample this:

$$v'(S_t) = \mathbf{R}_{t+1} + \gamma v(S_{t+1})$$

This is likely quite noisy — better to take a small step (with parameter  $\alpha$ ):

$$v'(S_t) = v(S_t) + \alpha[\mathbf{R}_{t+1} + \gamma v(S_{t+1}) - v(S_t)]$$

Remember in **Incremental Monte Carlo Updates**:

$$v'(S_t) \leftarrow v(S_t) + \alpha(\mathbf{G}_t - v(S_t))$$

This is likely quite noisy — better to take a small step (with parameter  $\alpha$ ):

$$v'(S_t) = v(S_t) + \alpha[R_{t+1} + \gamma v(S_{t+1}) - v(S_t)]$$

$R_{t+1} + \gamma v(S_{t+1})$  is called the **Temporal Difference Target**

$R_{t+1} + \gamma v(S_{t+1}) - v(S_t)$  is called the **Temporal Difference Error**

$\alpha$ , associated with the learning rate, can take several values

$R_{t+1} + \gamma v(S_{t+1})$  is called the **Temporal Difference Target**

Looking at the target, we notice that we are only taking into consideration the next immediate reward.

This is one specific form of Temporal Differences learning, namely TD(0).

For example, in TD(3):

$R_{t+1} + \gamma^1 R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v(S_{t+1})$  is the **Temporal Difference Target**

Note that TD( $\infty$ ), would mean that the target  $R_{t+1} + \gamma R_{t+2} + \dots = G_t$

**Bootstrapping:** update involves an estimate:

- Monte Carlo does not bootstrap.
- Dynamic Programming bootstraps.
- Temporal Differences learning bootstraps.

**Sampling:** update samples an expectation:

- Monte Carlo samples.
- Dynamic Programming does not sample.
- Temporal Differences learning samples.



We can apply the same idea to the  $Q$  function (action values)

Temporal-difference learning for action values:

- Update value  $q(S_t, A_t)$
- Towards estimated return  $R_{t+1} + \gamma \cdot q(S_{t+1}, A_{t+1})$

$$q'(S_t, A_t) \leftarrow q(S_t, A_t) + \alpha [R_{t+1} + \gamma \cdot q(S_{t+1}, A_{t+1}) - q(S_t, A_t)]$$

This algorithm is known as **SARSA**.

Temporal Difference is model-free (no knowledge of MDP) and learn directly from experience.

Temporal Difference can learn from incomplete episodes, by bootstrapping.

Temporal Difference can learn during each episode.

Temporal Difference can learn before knowing the final outcome:

- Temporal Difference can learn online after every step.
- Monte Carlo must wait until end of episode before return is known.

Temporal Difference can learn without the final outcome:

- Temporal Difference can learn from incomplete sequences.
- Monte Carlo can only learn from complete sequences.
- Temporal Difference works in continuing (non-terminating) environments.
- Monte Carlo only works for episodic (terminating) environments.

Temporal Difference is independent of the temporal span of the prediction:

- Temporal Difference can learn from single transitions.
- Monte Carlo must store all predictions (or states) to update at the end of an episode.

Temporal Difference needs reasonable value estimates.

Monte Carlo return  $G_t = R_{t+1} + \gamma \cdot R_{t+2} + \dots$  is an unbiased estimate of  $v_\pi(S_t)$ .

Temporal Difference target  $R_{t+1} + \gamma v(S_{t+1})$  is a biased estimate of  $v_\pi(S_t)$ .

unless  $v(S_{t+1}) = v_\pi(S_{t+1})$  where  $\pi$  is an optimal policy.

But the Temporal Difference target has lower variance:

- Return depends on many random actions, transitions, rewards.
- Temporal Difference target depends on one random action, transition, reward.

In some cases, Temporal Differences can have an irreducible bias.

The world may be partially observable:

- Monte Carlo would implicitly account for all the latent variables.

The function to approximate the values may fit poorly.

In the tabular case, both Monte Carlo and Temporal Differences will converge.



Temporal Differences **exploits Markov** property:

- Can help in **fully-observable** environments.

Monte Carlo **does not exploit Markov** property:

- Can help in **partially-observable** environments.

With finite data, or with function approximation, **the solutions may differ.**

Monte Carlo has **high variance, zero bias**:

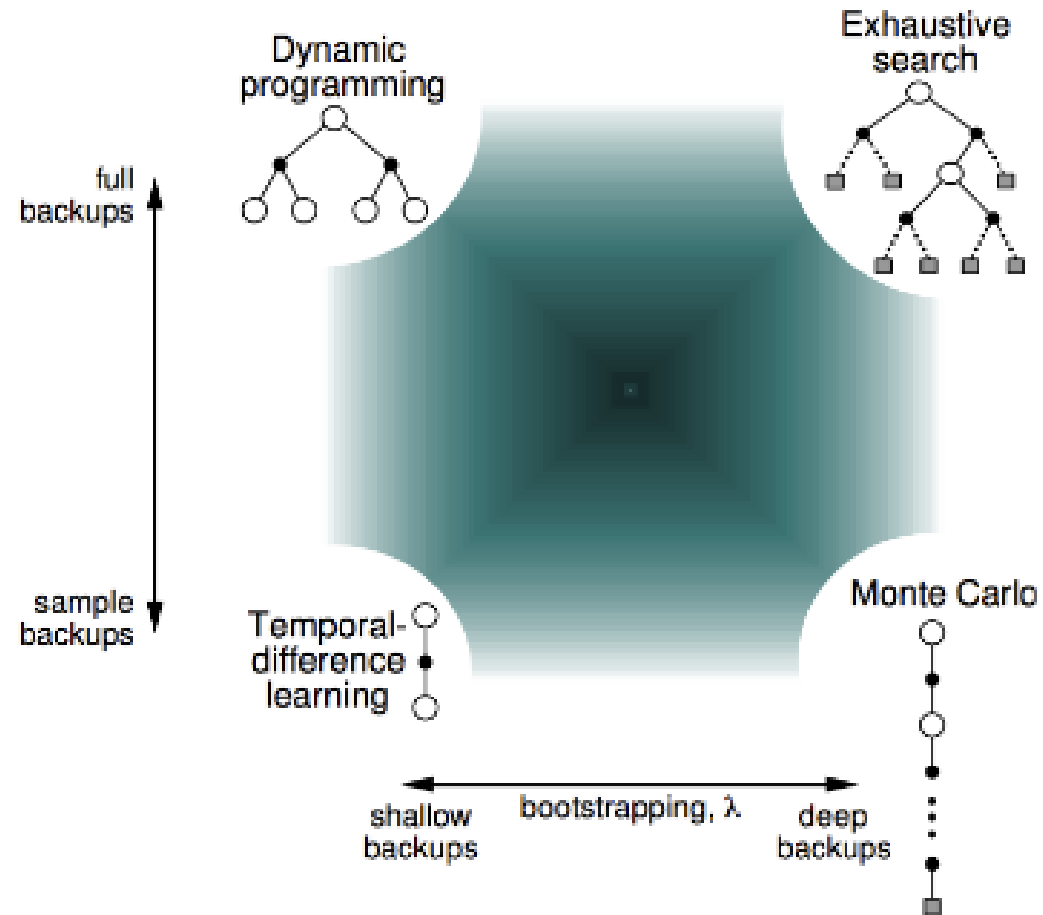
- Good convergence properties.
- Not very sensitive to initial value.
- Very simple to understand and use.

Temporal Differences **has low variance, some bias**:

- Usually more efficient than Monte Carlo.
- Temporal Differences usually converges to  $v_{\pi}(s)$ .
- More sensitive to initial value.

## Model Free Approaches

## Unified View of Reinforcement Learning



Temporal Differences uses value estimates which might be inaccurate:

- In addition, information can propagate back quite slowly.
- In Monte Carlo information propagates faster, but the updates are noisier.
- We can go in between Temporal Differences and Monte Carlo.

$R_{t+1} + \gamma v(S_{t+1})$  is called the **Temporal Difference Target**

Looking at the target we notice that we are only taking into consideration the next immediate reward.

This is one particular form of Temporal Differences learning, namely TD(0).

For example in TD(3):

$R_{t+1} + \gamma^1 R_{t+2} + \gamma^2 R_{t+3} + \gamma^3 v(S_{t+3})$  is the **Temporal Difference Target**


Note that TD( $\infty$ ), would mean that the target  $R_{t+1} + \gamma R_{t+2} + \dots = G_t$

Multi-step returns have benefits from both Temporal Differences and Monte Carlo:

- Bootstrapping can have issues with bias.
- Monte Carlo can have issues with variance.


## Visual Interpretation

$$v(S_t) \leftarrow \mathbb{E}[R_{t+1} + \gamma v(S_{t+1}) \mid A_t \sim \pi(S_t)]$$




Dynamic Programming

$$v(S_t) \leftarrow v(S_t) + \alpha (G_t - v(S_t))$$



Monte Carlo

$$v(S_t) \leftarrow v(S_t) + \alpha (R_{t+1} + \gamma v(S_{t+1}) - v(S_t))$$

Temporal Differences  
TD(0)

Let's discuss **control** methods!



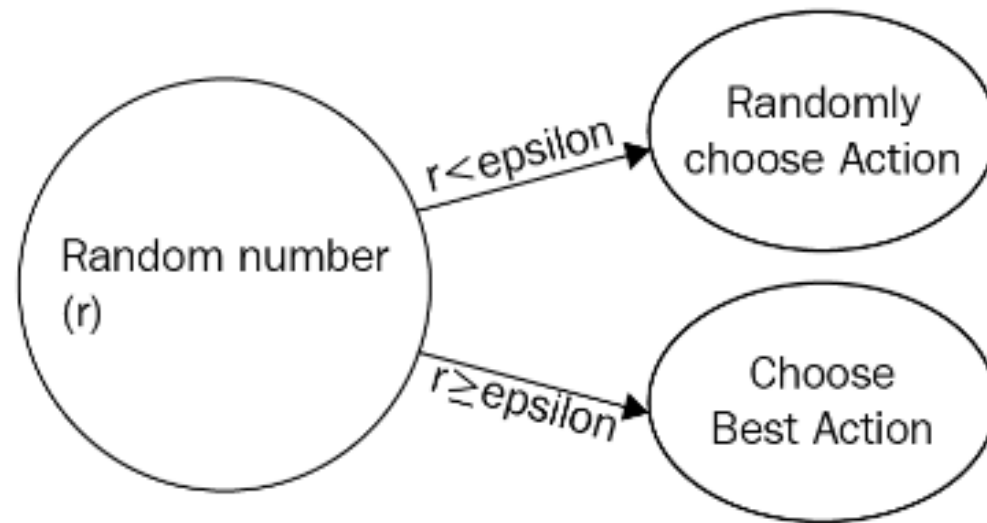
Remember two of the fundamental problems in reinforcement learning

- **Exploration** finds more information about the environment.
- **Exploitation** exploits known information to maximise reward.
- It is usually important to explore as well as exploit.

Simplest idea for ensuring continual exploration.

All actions are tried with non-zero probability.

With a given probability choose the greedy action.





# Model Free Control

# Model Free Policy Iteration

Policy Iteration is seen as a **greedy** approach.

The idea is to sequentially improve a value function (**policy evaluation**) and then a policy (**policy improvement**).



Greedy policy improvement over  $v(s)$  requires model of MDP:

$$\pi'(s) = \operatorname{argmax}_a E [\mathbf{R}_{t+1} + \gamma \cdot v(\mathbf{S}_{t+1}) \mid S_t = s, A_t = a]$$

Greedy policy improvement over  $q(s, a)$  is **model free**:


$$\pi'(s) = \operatorname{argmax}_a q(s, a)$$

This makes action values convenient.

Monte-Carlo policy evaluation,  $q \approx q_\pi$ .

Greedy policy improvement? No exploration!


(Can't sample all states and actions when learning by interacting)



Every episode:

Policy evaluation Monte-Carlo policy evaluation,  $q \approx q_\pi$ .

Policy improvement greedy policy improvement.



## Definition

Greedy in the Limit with Infinite Exploration (GLIE)

- All state-action pairs are explored infinitely many times:

$$\forall s, a \quad \lim_{t \rightarrow \infty} N(s, a) = \infty$$

- The policy converges to a greedy policy:

$$\lim_{t \rightarrow \infty} \pi_t(a|s) = I[a = \operatorname{argmax}_{a'} q(s, a')]$$


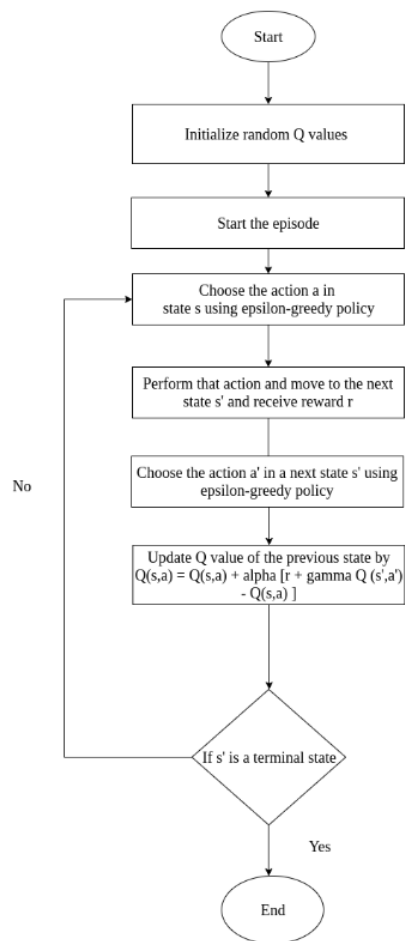
## Theorem

GLIE Model-free control **converges to the optimal action-value function**,  $q \rightarrow q_*$



# Model Free Control

## Updating Action-Value Functions with SARSA



Every time-step:

Policy evaluation SARSA,  $q \approx q_\pi$

Greedy policy improvement

$$q'(S_t, A_t) \leftarrow q(S_t, A_t) + \alpha [R_{t+1} + \gamma \cdot q(S_{t+1}, A_{t+1}) - q(S_t, A_t)]$$

## Theorem

Tabular SARSA converges to the optimal action-value function,  $q \rightarrow q_*$ , if the policy is GLIE.

## Dynamic programming

- Policy Iteration
- Value Iteration

## Monte Carlo

## Temporal Differences

- Temporal Differences (SARSA)
- Q-Learning

## On-policy learning

Learn about behaviour policy  $\pi$  from experience sampled from  $\pi$ .

## Off-policy learning

Learn about target policy  $\pi$  from experience sampled from  $\mu$ .

Evaluate target policy  $\pi(a|s)$  to compute  $v_\pi(s)$  or  $q_\pi(s, a)$ .

While using behaviour policy  $\mu(a|s)$  to generate actions.

## On-policy learning

- “Learn on the job”.

## Off-policy learning

- “Look over someone’s shoulder”.

Off-policy TD and Q-learning.

# Thank You!

**Morada:** Campus de Campolide, 1070-312 Lisboa, Portugal

**Tel:** +351 213 828 610 | **Fax:** +351 213 828 611

Acreditações e Certificações da NOVA IMS

Cofinanciado por



UNIGIS



Computing  
Accreditation  
Commission

